

NYU Politics Data Lab Workshop: Data Visualization with R and ggplot2

Pablo Barberá

Department of Politics
New York University

email: *pablo.barbera@nyu.edu*
twitter: @p_barbera

October 15, 2013

Data Visualization with R and ggplot2

Purpose of workshop: to introduce tools to generate elegant and effective plots for our academic research.

Why R?

- Mature, widely used, open-source, easily extensible (5K packages on CRAN repository)
- Object-oriented programming language.
- Many built-in basic and advanced statistical tools.

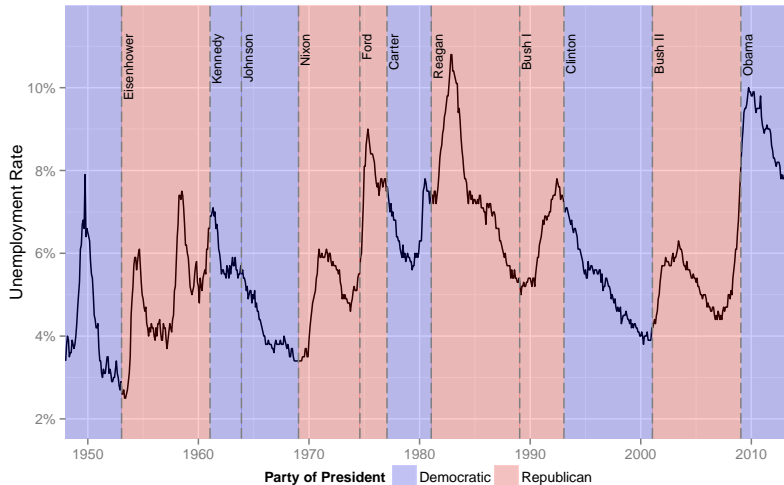
Why ggplot2?

- Based on “Grammar of Graphics” (Wilkinson, 2005)
 - powerful, consistent, modular.
- Sensible defaults, but also easy to customize
- Excellent online resources (and easy to google)


Outline

- ① Mastering the grammar of graphics
 - Building up a plot layer by layer
 - Scales, axes, legends
 - Themes and other options
- ② Applications
 - Annotated line plots
 - Regression coefficient plots
 - Network visualization
 - Maps and spatial analysis
 - Animated plots
- ③ Beyond ggplot2

Unemployment Rate in the United States, 1948–2013




What makes Facebook posts about politics popular?

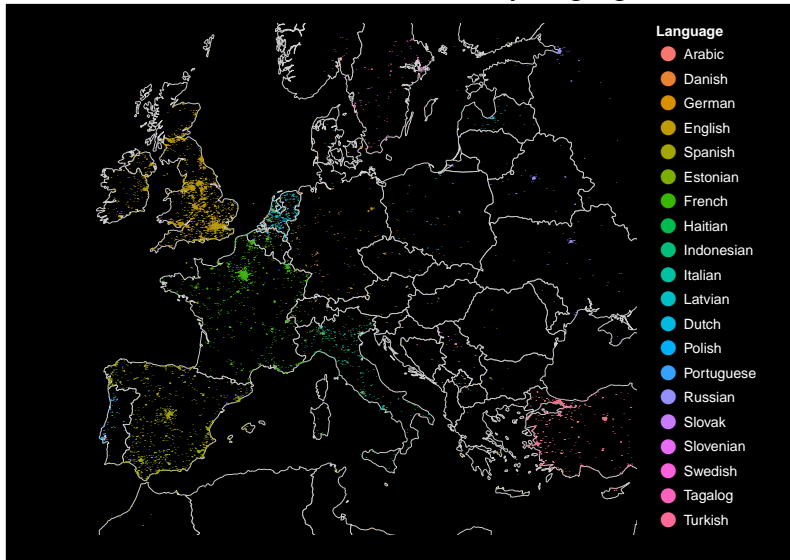


Data: 65K public Facebook posts about govt. shutdown

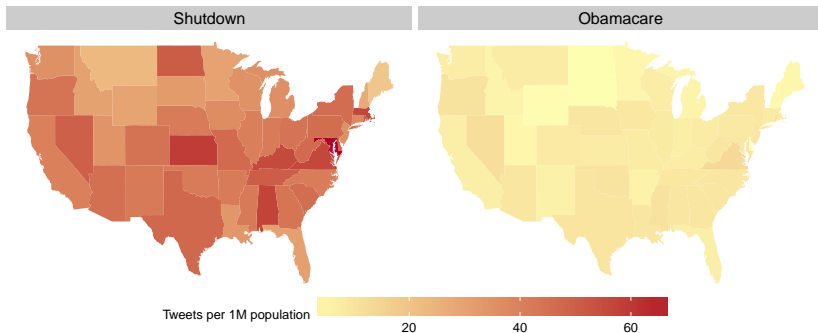
Visualizing your Facebook network



Geolocated tweets, colored by language



Geolocated tweets mentioning 'shutdown' and 'obamacare', 10/01



Animated plots showing convergence of MCMC algorithms

Materials

Code and data:

http://www.pablobarbera.com/workshop_ggplot2.zip

Short URL:

bit.ly/ggplot2_nyu

Fork my repo!

github.com/pablobarbera/Rdataviz

First R script (please run it now):

`code/00_installing_packages.R`

What is the grammar of graphics?

The grammar of graphics.

A statistical graph is a mapping from data to aesthetic attributes (color, shape, size) of geometric objects (points, lines, bars). The plot may also contain statistical transformations of the data and is drawn on a specific coordinate system. Faceting can be used to generate the same plot for different subsets of the data. It is the combination of these independent components that make up a graphic.



Hadley Wickham, *ggplot2*, page 3

What is the grammar of graphics?

Components of a graph:

- data** What you want to visualize, including variables to be mapped to aesthetic attributes.
- geom** Geometric objects representing what you see on the plot: points, lines, polygons, etc.
- stats** Statistical transformations of the data, such as binning or averaging.
- scales** Map values in the data space to values in an aesthetic space (color, shape, size...)
- coord** Coordinate system; provides axes and gridlines to make it possible to read the graph.
- facets** Breaking up the data into subsets, to be displayed independently on a grid

Resources and inspiration



Also: Stack Overflow, ggplot2 listserv & docs

Running example: analysis of Facebook data

Code: code/01_ggplot2-basics.R


Downloading public Facebook posts about the gov. shutdown.

```
library(devtools)
install_github("Rfacebook", "pablobarbera")
library(Rfacebook)
token <- "XXXXXXXXXX" # (your token here)
posts <- searchFacebook("shutdown", token, n=500)
users <- getUsers(posts$from_id, token)
names(users)[1] <- "from_id"
users <- users[!duplicated(users$from_id),]
fb.data <- merge(posts, users, by="from_id")
```

Univariate analysis: continuous variables

Distribution of number of likes for each post.


```
# loading library
library(ggplot2)
# base layer
p <- ggplot(fb.data,
  aes(x=likes_count))
# adding histogram
p + geom_histogram()
```



Univariate analysis: continuous variables

Distribution of (logged) number of likes for each post.

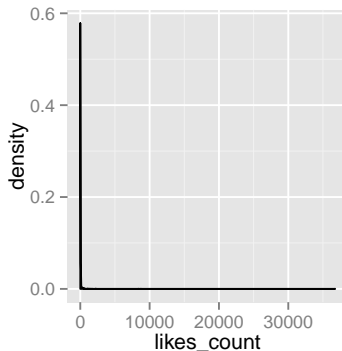
```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count))  
# adding histogram  
p + geom_histogram() +  
  scale_x_log10()
```



Univariate analysis: continuous variables

Distribution of number of likes for each post.

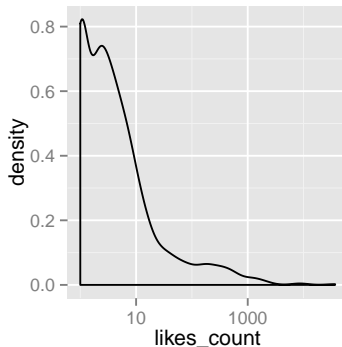
```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count))  
# adding density plot  
p + geom_density()
```



Univariate analysis: continuous variables

Distribution of (logged) number of likes for each post.


```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count))  
# adding density plot  
p + geom_density() +  
  scale_x_log10()
```



Univariate analysis: categorical variables

Distribution of posts by language


```
# first two characters of  
# 'locale' are lang. of user  
fb.data$language <- substr(  
  fb.data$locale, 1, 2)  
# base layer  
p <- ggplot(fb.data,  
  aes(x=language))  
# bar chart  
p + geom_bar()
```



Univariate analysis: categorical variables

Distribution of posts by language


```
# first two characters of  
# 'locale' are lang. of user  
fb.data$language <- substr(  
  fb.data$locale, 1, 2)  
# base layer  
p <- ggplot(fb.data,  
  aes(x=language))  
# bar chart (horizontal)  
p + geom_bar() +  
  coord_flip()
```



Bivariate analysis: two continuous variables

Counts of likes and counts of comments for posts


```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count,  
      y=comments_count))  
# scatter plot  
p + geom_point()
```



Bivariate analysis: two continuous variables

(Logged) counts of likes and (logged) counts of comments for posts


```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count,  
      y=comments_count))  
# scatter plot  
p + geom_point() +  
  scale_x_log10() +  
  scale_y_log10()
```



Bivariate analysis: two continuous variables

(Logged) counts of likes and (logged) counts of comments for posts


```
p <- ggplot(fb.data,  
  aes(x=likes_count,  
      y=comments_count))  
# scatter plot  
p + geom_point() +  
  scale_x_log10() +  
  scale_y_log10() +  
  stat_smooth(  
    na.rm=T,  
    data=fb.data[  
      fb.data$likes_count>0 &  
      fb.data$comments_count>0,]  
  )
```



Bivariate analysis: two continuous variables

Number of posts about shutdown, per hour

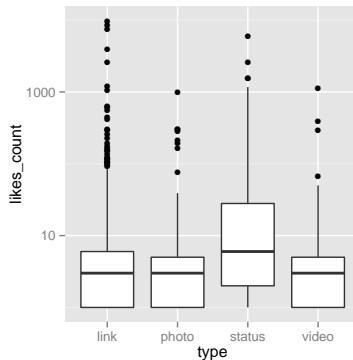
```
fb.data$time <- substr(fb.data$created_time, 1, 13) ## date + hour  
fb.data$time <- as.POSIXct(fb.data$time, format="%Y-%m-%dT%H")  
fb.data$count <- 1 ## counting number of posts per hour  
counts <- aggregate(fb.data$count, list(time=fb.data$time), sum)  
p <- ggplot(counts, aes(x=time, y=x))  
p + geom_line()
```



Bivariate analysis: continuous + categorical variables

Number of likes by type of post

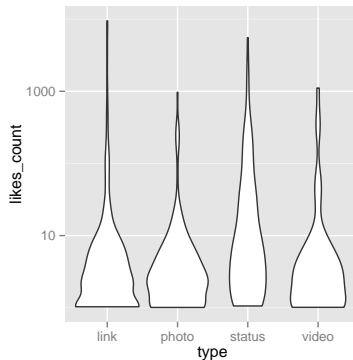
```
p <- ggplot(fb.data,  
  aes(x=type,  
      y=likes_count))  
# box plot  
p + geom_boxplot() +  
  scale_y_log10()
```



Bivariate analysis: continuous + categorical variables

Number of likes by type of post

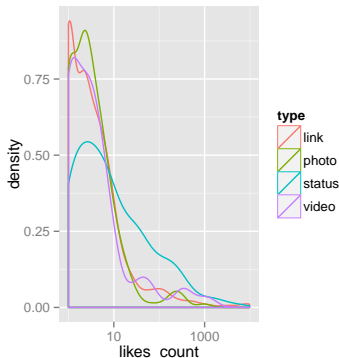
```
p <- ggplot(fb.data,  
  aes(x=type,  
      y=likes_count))  
# violin plot  
p + geom_violin() +  
  scale_y_log10()
```



Bivariate analysis: continuous + categorical variables

Number of likes by type of post

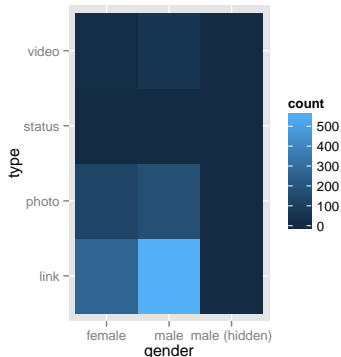
```
p <- ggplot(fb.data,  
  aes(x=likes_count))  
# density plot  
p + geom_density(  
  aes(color=type)) +  
  scale_x_log10()
```



Bivariate analysis: two categorical variables

Number of posts, by type and gender


```
# counts data.frame
tab <- data.frame(
  table(fb.data$gender,
        fb.data$type))
names(tab) <- c(
  "gender", "type", "count")
# base layer
p <- ggplot(tab,
  aes(x=gender, y=type))
# tile plot
p + geom_tile(
  aes(fill=count))
```



Multivariate analysis: three continuous variables

Number of likes and comments, by type of post

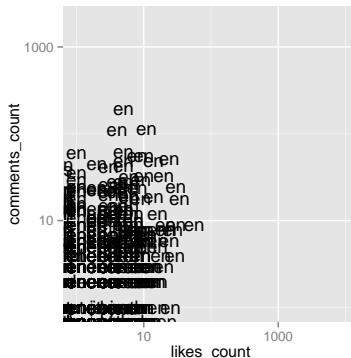
```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count,  
      y=comments_count))  
p + geom_point() +  
  scale_x_log10() +  
  scale_y_log10() +  
  facet_wrap(~type, nrow=2)
```



Multivariate analysis: continuous + categorical variables

Number of likes and comments, by country


```
# base layer
p <- ggplot(fb.data,
  aes(x=likes_count,
    y=comments_count,
    label=language))
p + geom_text() +
  scale_x_log10() +
  scale_y_log10()
```



Multivariate analysis: continuous + categorical variables

Number of posts per hour, by gender


```
counts.m <- aggregate(fb.data$count[fb.data$gender=="male"],  
  by=list(time=fb.data$time[fb.data$gender=="male"]), FUN=sum)  
counts.f <- aggregate(fb.data$count[fb.data$gender=="female"],  
  by=list(time=fb.data$time[fb.data$gender=="female"]), FUN=sum)  
counts.m$gender <- "male"; counts.f$gender <- "female"  
counts <- rbind(counts.m, counts.f)  
p <- ggplot(counts, aes(x=time, y=x, group=gender))  
p + geom_line(aes(color=gender))
```



Multivariate analysis: continuous + categorical variables

Number of likes and comments, by type of post

```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count,  
      y=comments_count))  
p + geom_point(  
  aes(color=type)) +  
  scale_x_log10() +  
  scale_y_log10()
```



Multivariate analysis: continuous + categorical variables

Number of likes and comments, by type of post

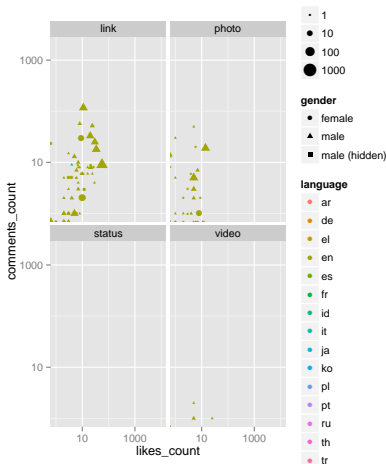
```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count,  
    y=comments_count))  
p + geom_point(  
  aes(color=type,  
    shape=type)) +  
  scale_x_log10() +  
  scale_y_log10()
```



Multivariate analysis: continuous + categorical variables

Number of likes and comments, with multiple scales


```
# base layer
p <- ggplot(fb.data,
  aes(x=likes_count,
    y=comments_count))
p + geom_point(
  aes(shape=gender,
    color=language,
    size=shares_count)) +
  scale_x_log10() +
  scale_y_log10() +
  scale_size(
    trans="log10") +
  facet_wrap(~type, nrow=2)
```



Dealing with overplotting

Counts of likes and counts of comments for posts

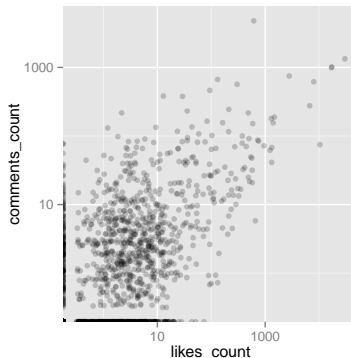
```
# base layer  
p <- ggplot(fb.data,  
  aes(x=likes_count,  
    y=comments_count))  
# jittered dots  
p + geom_jitter(  
  position = position_jitter(  
    width = .5, height=.5)) +  
  scale_x_log10() +  
  scale_y_log10()
```



Dealing with overplotting

Counts of likes and counts of comments for posts


```
# base layer
p <- ggplot(fb.data,
  aes(x=likes_count,
    y=comments_count))
# adding transparency
p + geom_jitter(
  position = position_jitter(
    width = .5, height=.5),
  alpha=1/5) +
  scale_x_log10() +
  scale_y_log10()
```



Dealing with overplotting

Counts of likes and counts of comments for posts

```
# base layer
p <- ggplot(fb.data[
  fb.data$likes_count>0 &
  fb.data$comments_count>0,],
  aes(x=likes_count,
      y=comments_count))
# hexagon binning
p + geom_hex() +
  scale_x_log10() +
  scale_y_log10() +
  scale_fill_continuous(
    trans="log10")
```




Customizing axes

Code: `code/02_scales_axes_legends.R`

Changing axis titles


```
p <- ggplot(fb.data,  
  aes(x=likes_count))  
p + geom_histogram() +  
  scale_x_continuous(  
    "Number of likes") +  
  scale_y_continuous(  
    "Post count")
```




Customizing axes

Changing axis limits (continuous variable)

```
p + geom_histogram() +  
  scale_x_continuous(  
    limits=c(0, 100))
```




```
p + geom_histogram() +  
  scale_x_log10(  
    limits=c(1, 100))
```



Customizing axes

Changing axis limits (categorical variable)


```
fb.data$language <- substr(  
  fb.data$locale, 1, 2)  
p <- ggplot(fb.data,  
  aes(x=language))  
p + geom_bar() +  
  scale_x_discrete(  
    limits=c("en", "es", "ru"))
```




Customizing axes

Changing axis breaks and labels (continuous variable)

```
p <- ggplot(fb.data,  
  aes(x=likes_count))  
p + geom_histogram() +  
  scale_x_continuous(  
    breaks=c(1, 10, 100, 1000))
```




```
p <- ggplot(fb.data,  
  aes(x=likes_count))  
p + geom_histogram() +  
  scale_x_log10(  
    breaks=c(1, 10, 100, 1000, 10000),  
    labels=c(1, 10, 100, "1K", "10K"))
```



Customizing axes


Changing axis limits (categorical variable)

```
p <- ggplot(fb.data,  
  aes(x=language))  
p + geom_bar() +  
  scale_x_discrete(  
    limits=c("en", "es", "ru"),  
    labels=c("en" = "English",  
             "es" = "Spanish",  
             "ru" = "Russian"))
```




Customizing axes

```
library(scales)
p <- ggplot(counts, aes(x=time, y=x))
p + geom_line() + scale_x_datetime(breaks="2 hours")
```




```
p + geom_line() + scale_x_datetime(breaks="2 hours",
  labels = date_format("%H"))
```




Customizing scales

Customizing properties of geoms

```
p <- ggplot(fb.data,  
  aes(x=likes_count))  
p + geom_histogram(  
  color="red")
```




```
p <- ggplot(fb.data,  
  aes(x=likes_count))  
p + geom_histogram(  
  fill="red")
```




Customizing scales


```
p <- ggplot(fb.data,  
  aes(x=likes_count,  
      y=comments_count)) +  
  scale_x_log10() +  
  scale_y_log10()  
p + geom_point( color="red")
```



```
p + geom_point(shape=15)
```




```
p + geom_point( size=1)
```



Customizing scales


Changing labels in legend

```
p + geom_point(  
  aes(shape=type)) +  
  scale_shape_manual(  
    "Post type",  
    limits = c("link", "status",  
               "video", "photo"),  
    labels = c("Link", "Status",  
               "Video", "Photo"),  
    values=c(1, 3, 4, 5))
```




Facets

```
p + geom_point() + facet_wrap(~type)
```




Facets

```
p + geom_point() +  
  facet_wrap(~type, nrow=4)
```




Facets

```
p + geom_point() + facet_wrap(~type, ncol=4)
```




Facets

```
p + geom_point() + facet_grid(gender~type)
```




Themes


```
p + geom_point() + theme_bw()
```




```
p + geom_point() + theme_grey()
```



```
p + geom_point() + theme_classic()
```




```
p + geom_point() + theme_minimal()
```



Saving plots

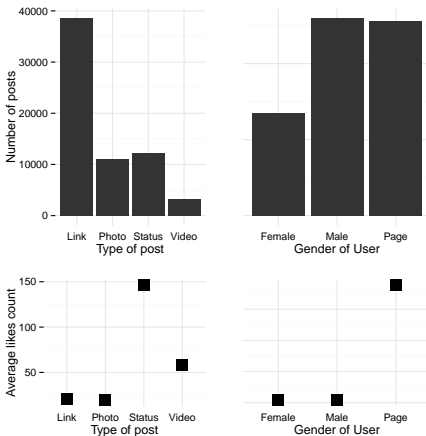
```
p + geom_point() + facet_grid(gender~type)
```



```
ggsave(pq, file="plots/plot.pdf", height=6, width=6)
```

Creating grid of plots

- 1 Store each plot as an object in memory
- 2 Use `pdf`, `grid`, `arrange` and `arrangeGrob` (from `gridExtra` package) to create plot.



Example: unemployment rate in the U.S.

Code: code/03_line_plots.R

```
library(gdata)
d <- read.xls("../data/unemployment_data.xls", stringsAsFactors=F)
# subsetting only what we need
months <- as.character(d[9, 2:13])
years <- as.character(d[10:75, 1])
d <- as.numeric(unlist(t(d[10:75, 2:13])))
# putting it together into a data frame
df <- data.frame(expand.grid(months, years))
names(df) <- c("month", "year")
df$unemp <- d/100
# removing missing values
df <- df[!is.na(df$unemp),]
```


Example: unemployment rate in the U.S.

```
# creating a date variable
df$date <- paste("01", df$month, df$year)
df$date <- as.Date(df$date, format="%d %b %Y")
# this is what the data looks like...
str(df)

## 'data.frame': 788 obs. of 4 variables:
## $ month: Factor w/ 12 levels "Jan","Feb","Mar",...: 1 2 3 4 5 6 7 8
## $ year : Factor w/ 66 levels "1948","1949",...: 1 1 1 1 1 1 1 1 1 1
## $ unemp: num 0.034 0.038 0.04 0.039 0.035 0.036 0.036 0.039 0.038
## $ date : Date, format: "1948-01-01" "1948-02-01" ...
```


Example: unemployment rate in the U.S.

```
# plot with default options  
library(ggplot2)  
p <- ggplot(df, aes(x=date, y=unemp))  
p + geom_line()
```




Example: unemployment rate in the U.S.


```
p + geom_line(color='red')
```




```
p + geom_line(size=10)
```



```
p + geom_line(linetype=3)
```




```
p + geom_line(size=0.1)
```




Example: unemployment rate in the U.S.

```
library(scales)
# customizing scale
p + geom_line() + ## full date
  scale_x_date(labels = date_format("%Y/%m/%d"))
```




Example: unemployment rate in the U.S.

```
# customizing scale  
p + geom_line() + ## full month and year  
  scale_x_date(labels = date_format("%B %Y"))
```




Example: unemployment rate in the U.S.

```
# customizing scale  
p + geom_line() + ## every 10 years  
  scale_x_date(labels = date_format("%Y"),  
    breaks = date_breaks("10 years"))
```




Example: unemployment rate in the U.S.

```
p + geom_line() + ## manual breaks  
scale_x_date(labels = date_format("%Y"),  
  breaks = as.Date(c("1950-01-01", "1965-01-01", "1980-01-01",  
    "1980-01-01", "1995-01-01", "2010-01-01")))
```



Example: unemployment rate in the U.S.

```
p + geom_line() + ## customizing y scale  
scale_y_continuous("Unemployment Rate", labels = percent,  
  limits=c(0.02, 0.115), breaks=c(0.02, 0.04, 0.06, 0.08, 0.10))
```




Example: unemployment rate in the U.S.


```
p + geom_line() + theme_bw()
```



```
p + geom_line() + theme_grey()
```



```
p + geom_line() + theme_classic()
```



```
p + geom_line() + theme_minimal()
```




Example: unemployment rate in the U.S.

```
# background colors for different presidents
name <- c("Truman", "Eisenhower", "Kennedy", "Johnson", "Nixon",
          "Ford", "Carter", "Reagan", "Bush I", "Clinton",
          "Bush II", "Obama")
start <- as.Date(c("1945-04-12", "1953-01-20", "1961-01-20",
                  "1963-11-22", "1969-01-20", "1974-08-09", "1977-01-20",
                  "1981-01-20", "1989-01-20", "1993-01-20", "2001-01-20",
                  "2009-01-20"))
end <- c(start[-1], as.Date("2013-10-15"))
party <- c("D", "R", "D", "D", "R", "R", "D", "R", "R", "D",
          "R", "D")
pres <- data.frame(name, start, end, party, stringsAsFactors=F)
pq <- p + geom_line() + scale_y_continuous("Unemployment Rate",
      labels = percent, limits = c(0.02, 0.115),
      breaks = c(0.02, 0.04, 0.06, 0.08, 0.10))
```



Example: unemployment rate in the U.S.

```
yrng <- range(df$unemp, na.rm=TRUE)
xrng <- range(df$date, na.rm=TRUE)
pq + geom_rect(data=pres, aes(NULL, NULL, xmin=start, xmax=end,
  fill=party), ymin=yrng[1]-0.05, ymax=yrng[2]+0.05) +
  scale_fill_manual(values = alpha(c("blue", "red"), 0.2))
```



Example: unemployment rate in the U.S.

```
pq + geom_rect(data=pres, aes(NULL, NULL, xmin=start, xmax=end,  
  fill=party), ymin=yrng[1]-0.05, ymax=yrng[2]+0.05) +  
  scale_fill_manual(values = alpha(c("blue", "red"), 0.2)) +  
  geom_text(data=pres, aes(x=start, y=yrng[2], label=name),  
  size=3, hjust=1, vjust=1.25, angle=90)
```




Example: unemployment rate in the U.S.

```
# final touches
library(grid)
yrng[2] <- yrng[2] + 0.05 # moving names up
pq <- pq + geom_rect(data=pres, aes(NULL, NULL, xmin=start,
  xmax=end, fill=party), ymin=yrng[1]-0.05, ymax=yrng[2]) +
scale_fill_manual(values = alpha(c("blue", "red"), 0.2)) +
scale_y_continuous("Unemployment Rate", labels = percent,
  limits = c(.02, .115), breaks=c(.02, .04, .06, .08, .10)) +
geom_text(data=pres, aes(x=start, y=yrng[2], label=name),
  size=3, hjust=1, vjust=1.25, angle=90) +
geom_vline(data=pres, aes(xintercept=as.numeric(start)),
  color="grey50", linetype=5, size=0.5) + # separation lines
theme(axis.title.x = element_blank(), # removing x title
  legend.position = "bottom", # legend to bottom
  legend.margin = unit(-1, "cm") ) + # less margin
scale_fill_manual("Party of President", # fixing labels
  values=alpha(c("blue", "red"), 0.2),
  labels= c("Democratic", "Republican"))
```

Example: unemployment rate in the U.S.

```
pq
```



Example: determinants of popularity on Facebook

Code: code/04_coefficients_plots.R

```
library(Rfacebook) # downloading public FB. posts
token <- "XXXXXXXXXX" # (your token here)
fb.data <- searchFacebook("shutdown", token, n=2000)
# recoding data
fb.data$gender[is.na(fb.data$gender)] <- "Page"
fb.data$gender[fb.data$gender == "male (hidden)"] <- "male"
fb.data$language <- substr(fb.data$locale, 1, 2)
fb.data$english <- ifelse(
  fb.data$language == "en" & !is.na(fb.data$language),
  "English", "Others")
```

Example: determinants of popularity on Facebook

```
fb.data$obamacare <- grepl("obamacare", fb.data$message,  
  ignore.case=TRUE)  
fb.data$boehner <- grepl("boehner", fb.data$message,  
  ignore.case=TRUE)  
fb.data$furlough <- grepl("furlough", fb.data$message,  
  ignore.case=TRUE)  
fb.data$time <- substr(fb.data$created_time, 12, 13)  
fb.data$night <- fb.data$time %in%  
  c("00", "01", "02", "03", "04", "05")  
fb.data$morning <- fb.data$time %in%  
  c("06", "07", "08", "09", "10", "11")  
fb.data$afternoon <- fb.data$time %in% as.character(12:17)  
fb.data$evening <- fb.data$time %in% as.character(18:23)
```

Example: determinants of popularity on Facebook

```
# Running regressions
```

```
r1 <- summary(lm(log(likes_count+1) ~ gender + type + english +  
  obamacare + boehner + furlough + morning + afternoon +  
  evening, data=fb.data))  
r2 <- summary(lm(log(comments_count+1) ~ gender + type +  
  english + obamacare + boehner + furlough + morning +  
  afternoon + evening, data=fb.data))  
r3 <- summary(lm(log(shares_count+1) ~ gender + type +  
  english + obamacare + boehner + furlough + morning +  
  afternoon + evening, data=fb.data))
```

Example: determinants of popularity on Facebook


```
# Saving regressions into a data frame
df <- lapply(list(r1, r2, r3), function(x)
  data.frame(
    var = rownames(x$coefficients)[2:13],
    coef = x$coefficients[2:13, "Estimate"],
    sd = x$coefficients[2:13, "Std. Error"])
)
df <- do.call(rbind, df)
df$dv <- rep(c("likes count", "comments count",
  "shares count"), each=12)
```


Example: determinants of popularity on Facebook

```
# changing variable labels
levels(df$var) <- c("Posted in afternoon", 'Mentions "boehner"',
  "Post in English", "Posted in evening", 'Mentions "furlough*"',
  "Author is male", "Author is page", "Posted in morning",
  'Mentions "obamacare"', "Post is a photo",
  "Post is a status update", "Post is a video")
df$var <- factor(df$var, levels=c("Author is male", "Author is page",
  "Posted in morning", "Posted in afternoon", "Posted in evening",
  "Post in English", 'Mentions "boehner"', 'Mentions "furlough*"',
  'Mentions "obamacare"', "Post is a photo",
  "Post is a status update", "Post is a video"))
```

Example: determinants of popularity on Facebook


```
library(ggplot2)
p <- ggplot(df, aes(y=coef, x=var))
p + geom_point(aes(color=dv)) + coord_flip()
```



Example: determinants of popularity on Facebook


```
# adding error bars
```

```
p + geom_point(aes(color=dv)) + coord_flip() +  
  geom_errorbar(aes(x=var, ymin=coef-2*sd, ymax=coef+2*sd))
```




Example: determinants of popularity on Facebook

```
# adding lines to indicate uncertainty  
p + geom_point(aes(color=dv)) + coord_flip() +  
  geom_linerange(aes(x=var, ymin=coef-2*sd, ymax=coef+2*sd))
```



Example: determinants of popularity on Facebook

```
# order matters! also use position_dodge to avoid overlap
p <- ggplot(df, aes(y=coef, x=var, group=dv))
p + coord_flip() + geom_linerange(aes(ymin=coef-2*sd,
  ymax=coef+2*sd, color=dv), position=position_dodge(.5)) +
  geom_point(aes(color=dv), position = position_dodge(.5))
```



Example: determinants of popularity on Facebook

```
# fixing axes, legends, line at 0... and putting it all together
library(scales)
p <- ggplot(df, aes(y=coef, x=var))
pq <- p + coord_flip() +
  # adding first (long, thin) line for coef +/- 2 sd
  geom_linerange(aes(ymin=coef-2 *sd, ymax=coef+2*sd, color=dv),
    position=position_dodge(.5)) +
  # adding second (short, thick) line for coef +/- 1 sd
  geom_linerange(aes(ymin=coef-sd, ymax=coef+sd, color=dv),
    position=position_dodge(.5), size=1) +
  # changing y axis title and scale (y and not bc of coord_flip)
  scale_y_continuous("% increase in popularity metric",
    labels = percent) +
  # adding line at 0
  geom_hline(yintercept=0, linetype=5, color="grey50")
```

Example: determinants of popularity on Facebook


```
# changing theme and removing axis title
pq <- pq + theme_bw() + theme(
  axis.title.y=element_blank() ) +
  # changing title of legend
  scale_color_discrete("Effect on...") +
  # horizontal lines to separate different types of variables
  geom_vline(xintercept=2.5, color="grey60") +
  geom_vline(xintercept=5.5, color="grey60") +
  geom_vline(xintercept=6.5, color="grey60") +
  geom_vline(xintercept=9.5, color="grey60")
```

Example: determinants of popularity on Facebook

```
# adding labels indicating types of variables
pq <- pq + annotate("text", x=9.8, y=-.32, label="Format",
  size=3, color="grey60", hjust=0) +
  annotate("text", x=6.8, y=-.32, label="Content",
  size=2.8, color="grey60", hjust=0) +
  annotate("text", x=5.8, y=-.32, label="Language",
  size=2.8, color="grey60", hjust=0) +
  annotate("text", x=2.8, y=-.32, label="Time",
  size=2.8, color="grey60", hjust=0) +
  annotate("text", x=0.8, y=-.32, label="Author",
  size=2.8, color="grey60", hjust=0)
```


Example: determinants of popularity on Facebook

pq




Example: visualizing your Facebook network

Code: code/05_networks.R

```
library(Rfacebook) # downloading network of FB friends
token <- "XXXXXXXXXX" # (your token here)
mat <- getNetwork(token, format="adj.matrix")
# preparing node list and layout with igraph
library(igraph)
network <- graph.adjacency(mat, mode="undirected")
fc <- fastgreedy.community(network) ## communities
l <- layout.fruchterman.reingold(network, niter=1000, coolexp=0.5)
# preparing data for plot
d <- data.frame(l); names(d) <- c("x", "y")
d$cluster <- factor(fc$membership)
```


Example: visualizing your Facebook network

```
p <- ggplot(d, aes(x=x, y=y, color=cluster))  
(pq <- p + geom_point())
```



Example: visualizing your Facebook network

```
## too many clusters! let's pick just those with 10 friends or more  
large.clusters <- which(table(fc$membership)>=10)  
fc$membership[fc$membership %in% large.clusters == FALSE] <- "Others"  
d$cluster <- factor(fc$membership)  
p <- ggplot(d, aes(x=x, y=y, color=cluster))  
(pq <- p + geom_point())
```




Example: visualizing your Facebook network

```
## simplify even further by keeping only nodes in giant component
cl <- clusters(network)
gc <- which(cl$membership == 1)
mat <- mat[gc, gc]
network <- graph.adjacency(mat, mode="undirected")
fc <- fastgreedy.community(network)
set.seed(123)
l <- layout.fruchterman.reingold(network, niter=1000, coolexp=0.5)
d <- data.frame(l); names(d) <- c("x", "y")
d$cluster <- factor(fc$membership)
```


Example: visualizing your Facebook network

```
p <- ggplot(d, aes(x=x, y=y, color=cluster))  
(pq <- p + geom_point())
```




Example: visualizing your Facebook network

```
## now let's add the edges
edgelist <- get.edgelist(network, names=FALSE)
edges <- data.frame(d[edgelist[,1],c("x", "y")],
                    d[edgelist[,2],c("x", "y")])
names(edges) <- c("x1", "y1", "x2", "y2")
(pq <- pq + geom_segment(
  aes(x=x1, y=y1, xend=x2, yend=y2),
  data=edges, size=0.25, color="grey", alpha=1/3))
```




Example: visualizing your Facebook network

```
## note that the order matters!  
p <- ggplot(d, aes(x=x, y=y, color=cluster))  
(pq <- p + geom_segment(aes(x=x1, y=y1, xend=x2,  
  yend=y2), data=edges, size=0.25, color="grey",  
  alpha=1/3) + geom_point())
```




Example: visualizing your Facebook network

```
## change a few of the theme options to make it look better
(pq <- pq + theme(
  panel.background = element_rect(fill = "black"),
  plot.background = element_rect(fill="black"),
  axis.line = element_blank(), axis.text = element_blank(),
  axis.ticks = element_blank(), axis.title = element_blank(),
  panel.border = element_blank(), panel.grid.major =
    element_blank(), panel.grid.minor = element_blank()))
```




Example: visualizing your Facebook network

```
## now let's customize the legend
(pq <- pq + theme(legend.background =
  element_rect(colour = "white", fill = "black"),
  legend.key = element_rect(fill = "black", colour = F),
  legend.title = element_text(color="white"),
  legend.text = element_text(color="white")))
```



Example: visualizing your Facebook network

```
## after exploring who is in each cluster,  
## labeling each of them by its category  
labels <- c("NYU", "EUI", "UPF", "MA", "Caixa", "NYC", "Others")  
(pq <- pq + scale_color_discrete(labels=labels))
```




Example: visualizing your Facebook network



```
## putting it all together...
p <- ggplot(d, aes(x=x, y=y, color=cluster))
pq <- p + geom_segment(
  aes(x=x1, y=y1, xend=x2, yend=y2),
  data=edges, size=0.25, color="white", alpha=1/3) +
  ## note that here I add a border to the points
  geom_point(color="grey20", aes(fill=cluster), shape=21, size=2) +
  scale_fill_discrete(labels=labels) +
  theme(
    panel.background = element_rect(fill = "black"),
    plot.background = element_rect(fill="black"),
    axis.line = element_blank(), axis.text = element_blank(),
    axis.ticks = element_blank(),
    axis.title = element_blank(), panel.border = element_blank(),
    panel.grid.major = element_blank(),
    panel.grid.minor = element_blank(),
    legend.background = element_rect(colour = F, fill = "black"),
    legend.key = element_rect(fill = "black", colour = F),
    legend.title = element_text(color="white"),
    legend.text = element_text(color="white")
  ) +
  ## changing size of points in legend
  guides(fill = guide_legend(override.aes = list(size=5)))
```

Example: visualizing your Facebook network

pq




Beyond ggplot2




Data-Driven Documents (D3), in javascript

Beyond ggplot2




Dygraphs (javascript)

Beyond ggplot2



googleVis for R



rCharts

NYU Politics Data Lab Workshop: Data Visualization with R and ggplot2

Pablo Barberá

Department of Politics
New York University

email: *pablo.barbera@nyu.edu*
twitter: @p_barbera

October 15, 2013